

Sleep Soundly: Reliable Alerting with Unit Testing in Prometheus


Rubén Cougil Grande, DevBcn 2023

Being on-call

- **Availability** outside working hours (nights, weekends, holidays...).
- 🔥 **False Positives** due to flawed alerting rules.
- Needs of **Real Data**, hard to test these rules in advance.






Unit Testing in Prometheus


-  built-in cli tool: **promtool**
- **Assert** the alerting rule perform as intended.
- **Early detection** of issues.
- Integrate it in **CI** pipeline.
- == sleep better 😴



How Alerting Rules works in Prometheus

- Evaluation of expressions written in **PromQL**
 - `sum(up{job="app"}) == 0`
- **States:**  Inactive  Pending  Firing.
- Custom **Labels**, such as “severity”.
- **Annotations**, such as “summary” or “runbook URL”.
- 🤖 **PromQL** can be a tricky DSL.

PromQL complexity

-  The complexity of the query increases the chances of **errors**.
- **Unit Tests** will help.

```
sum by(namespace, application) (increase(http_server_requests_seconds_count{namespace="live",  
application="backend", status=~"500|503"}[10m]))
```

/


```
sum by(namespace, application) (increase(http_server_requests_seconds_count{namespace="live",  
application="backend"}[10m])) * 100 > 10
```

and ON()

```
sum by(namespace, application) (increase(http_server_requests_seconds_count{namespace="live",  
application="backend"}[10m])) > 100
```

How Alerting Rules works in Prometheus

```
- alert: InstancesDownV1
  expr: sum(up{job="app"}) == 0
  labels:
    severity: sev1
  annotations:
    summary: "All instances of App are down"
    description: "All instances of App are down"
```



How to unit test a Rule

- 🙌 Generate data for the Test Case (**Arrange**):

interval: 1m

input_series:

- **series:** 'up{job="app", instance="app-1:2223"}'
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
values: "0x14"
- **series:** 'up{job="app", instance="app-2:2223"}'
1 1 1 1 1 0 0 0 0 0 0 0 0 0 1 1 1 1
values: "1x4 0x9 1x4"

- 🙌 Set the expectations (**Assert**):

- eval_time: 4m
alertname: InstancesDownV1
- eval_time: 5m
alertname: InstancesDownV1
exp_alerts:
 - exp_labels:
severity: page
exp_annotations:
summary: "All instances of the App are down"
description: "All instances of the App are down"
- eval_time: 15m
alertname: InstancesDownV1

- 🙌 Run the test (using **promtool**):
 - `$> promtool test rules instances_down_v1_test.yml`
SUCCESS

- ! How output would look like **if assertion fails**:

- `$> promtool test rules instances_down_v1_test.yml`

FAILED:

alertname: InstancesDownV1, time: **4m**,

exp: [

Labels: {alertname="InstancesDownV1", severity="page"}

Annotations: {description="All instances of the App are down", summary="All instances of the App are down"}

],

got: []

Samples can be empty (“_”) or stale

- In previous example, we’ve set “0” when target is down in the test.
- In real world, target down will report **empty sample** (“_”)
- More accurate example:

input_series:

- series: 'up{job="app", instance="app-1:2223"}'
_ _ _ _ _ _ _ _ _ _
values: "_x14"
- series: 'up{job="app", instance="app-2:2223"}'
1 1 1 1 1 _ _ _ _ _ _ _ _ 1 1 1 1 1
values: "1x4 _x9 1x4"

Samples can be empty (“_”) or stale

- 😬 Now the **test fails**, because:
 - `expr`: `sum(up{job="app"}) == 0`
 - Sum of empty state sample (NaN) will never be equals to zero.
- 🙌 Expression should be:
 - `expr`: `sum(up{job="app"} OR on() vector(0)) == 0`
 - `vector(0)` returns zero when sample is empty.
 - ✅ Now the test passes

Usage of “for” in rules






- **States:** ● Inactive ● Pending ● Firing.
 - ● Pending state depends on “**for**” clause
 - **!** It will **affect** test **expectations**
-
- `alert: InstancesDownV1`
 - `expr: sum(up{job="app"}) == 0`
 - `for: 5m` ←
 - `labels:`
 - `severity: sev1`
 - `annotations:`
 - `summary: "All instances of App are down"`
 - `description: "All instances of App are down"`

Scalability and Alerting Rules

- 🙄 Prometheus server **does NOT scale horizontally**
- Configure Prometheus to send metrics to an **external service** (*Thanos, Cortex...*)
- Alerting Rules are evaluated using Prometheus internal TSDB.
- 15 days of **retention** by default.
- 🙅 Do not create rules that go beyond that.
- Keep this in mind in tests.



Benefits of Unit Testing Rules

-  Ensuring rule correctness.
-  Early detection of issues.
-  Facilitating code refactoring.
-  Supporting continuous integration.
-  Collaboration and documentation.

Resources

- Github Repository with examples:
 - <https://github.com/rubencougil/prometheus-testing>
- Unit Testing in Prometheus official [page](#).
- Twitter: [@rcougil](#)
- Talk resources @ [Whova](#)
- [Awesome](#) Prometheus Alerts in Github

